

Ligax AI Hairstyle Classification Benchmark

Technical Documentation — Version 1.0

Maintained by: Ligax AI

Contact: admin@ligaxai.com

Benchmark Repository: https://github.com/Ligax-LTD/Lig_benchmark_test

Playground: <https://www.ligaxai.com/playground>

Dataset: <https://universe.roboflow.com/testing-workspace-htshl/woman-hairstyles-bevyz>

Table of Contents

Technical Documentation — Version 1.0	1
Table of Contents	1
1. Overview	3
2. Background and Motivation	3
2.1 Why a Public Benchmark	3
2.2 Why Afro-Caribbean Protective Hairstyles	3
3. Lig Vision: The Model Under Evaluation	4
4. Benchmark Taxonomy	4
4.1 Label Dimensions	4
4.2 Multi-Label Structure	5
4.3 Relationship to Model Internal Taxonomy	5
5. Dataset	5
5.1 Overview	5
5.2 Source and Construction	5
5.3 Labelling Methodology	6
5.4 Verification	6
5.5 Label Distribution	6
5.6 Known Label Convention Notes	6
6. Comparison Models	6
6.1 GPT-5.5	6

6.2 Claude Opus 4.8	6
6.3 Gemini 2.5 Flash	6
6.4 Prompting Approach	6
7. Evaluation Methodology	7
7.1 Head-Split Evaluation	7
7.2 Metrics	7
7.3 Threshold Handling	7
7.4 Hierarchical Confidence Boosting	7
7.5 Reproducibility	8
8. Current Results	8
8.1 Named Styles (Key) and Base Technique Categories	8
8.2 Named Styles -- Per-Class Detail (Lig Vision)	8
8.3 Base Technique Categories -- Per-Class Detail (Lig Vision)	9
8.4 Additional Dimensions	9
9. Reproducing Results	9
9.1 Requirements	9
9.2 Configuration	9
9.3 Running the Benchmark	10
9.4 Verifying Lig Vision Results	10
10. Limitations and Known Issues	10
11. Ongoing Development	11
12. Version History	11

1. Overview

This document describes the Ligax AI public hairstyle classification benchmark: its scope, design rationale, dataset, evaluation methodology, and results. The benchmark is designed to be fully reproducible by external parties using publicly available tools and datasets.

The benchmark evaluates Lig Vision, the hairstyle classification model powering the Ligax AI platform, alongside three frontier general-purpose vision-language models. It measures classification performance across multiple label dimensions simultaneously: specific named styles, base technique categories, technique modifiers, size and length attributes, and add-on or accessory labels.

All benchmark code, the dataset, and pre-computed model results are publicly available at https://github.com/Ligax-LTD/Lig_benchmark_test. Results for Lig Vision are independently verifiable image by image via the Ligax AI playground at <https://www.ligaxai.com/playground>.

2. Background and Motivation

2.1 Why a Public Benchmark

Hairstyle classification for booking applications requires a level of precision that general-purpose image recognition does not provide. Correctly identifying not just a broad style category but the specific technique, installation variant, size, and add-ons from a single image is a hard multi-label classification problem, particularly in the domain of protective and textured hairstyles where many styles are visually similar and the distinguishing features are technique-dependent.

Ligax AI publishes this benchmark for two reasons. The first is transparency: the benchmark is designed to give anyone the means to independently verify the classification accuracy we report. The second is rigour: comparing performance against frontier general-purpose models on a demanding domain-specific benchmark is a meaningful test of what specialised training contributes. A benchmark where generalist models easily match a fine-tuned specialist does not tell you much. A benchmark where the specialist leads substantially tells you something real about the value of domain-specific training data and label quality.

This is intended as a directional public benchmark. Ligax AI conducts more extensive internal evaluation across a broader range of styles and images. The public benchmark is a reproducible subset designed for external verification, not a complete picture of model capability.

2.2 Why Afro-Caribbean Protective Hairstyles

The benchmark focuses on Afro-Caribbean protective hairstyles as its initial evaluation domain. This was a deliberate choice made for several reasons.

Lig Vision is trained across a wide range of hairstyle categories spanning multiple regions, cultural traditions, and demographic groups. The model recognises hairstyles from diverse geographic and cultural contexts including West African, East African, Caribbean, South Asian, East Asian, and European styling traditions, among others. It is not a model built for a single demographic.

The decision to centre the initial public benchmark on Afro-Caribbean protective styles reflects the difficulty of this specific category, not a limitation of the model's scope. These styles present a classification challenge that is technically harder than most other hairstyle categories for several compounding reasons.

Visual similarity between distinct techniques. Many protective styles share the same silhouette, length, and colour profile. The features that distinguish one named style from another are often subtle: the texture of the extension material, the geometry of the base pattern, the strand thickness, the tip finish. A model that relies primarily on silhouette or colour recognition will fail systematically to resolve these distinctions. Classifying them correctly requires the model to have learned the domain at a level of granularity that general-purpose training does not provide.

Multi-label co-occurrence structure. A single image routinely carries multiple simultaneously correct labels across different levels of the label hierarchy. The specific named style, the base installation technique, the size attribute, and any add-on elements are all simultaneously correct and all independently relevant to the booking. Missing any label in the set is an operational error. This is a harder task than single-label or even standard multi-class classification.

Underrepresentation in general-purpose training corpora. The fine-grained variation within braided, twisted, and loc-based style categories is not well-represented in the training data underlying standard foundation vision models. A model that performs well on this benchmark can be expected to generalise effectively to less challenging hairstyle categories. The converse is not guaranteed. Starting the benchmark at the hardest end of the difficulty range makes the results maximally informative.

Subsequent benchmark versions will expand to additional style categories and regions. This initial version establishes the evaluation methodology and provides a baseline against which future improvements will be measured.

3. Lig Vision: The Model Under Evaluation

Lig Vision is the hairstyle classification model deployed in production within the Ligax AI platform. It is a fine-tuned deep learning model trained on a curated hairstyle image dataset with labels produced by professional hairstylists with active domain expertise.

The model is trained on a substantially broader taxonomy than this benchmark covers. It is capable of classifying hairstyles across many cultural and demographic contexts, recognising styles from multiple regions and traditions at multiple levels of specificity. The benchmark taxonomy is a subset selected for this evaluation domain.

The model produces simultaneous predictions across multiple label dimensions rather than a single class label. For any input image, it outputs probabilities across all label dimensions it was trained on, allowing the system to produce a complete descriptive profile of the hairstyle rather than a single flat classification.

The model is deployed in production at Ligax AI and processes real client uploads daily. The benchmark evaluation uses the same inference endpoint that powers the production system. There is no evaluation-specific model version. Results obtained via the playground at <https://www.ligaxai.com/playground> are produced by the same model that generated the benchmark results.

Reproducibility note: The Lig Vision inference endpoint is not published in the benchmark repository. Pre-computed results for all benchmark images are included in the repository at `results/lig\_vision.json`. Individual results can be verified via the Ligax AI playground.

4. Benchmark Taxonomy

The benchmark taxonomy defines the label space against which all models are evaluated. It is organised into five label dimensions, each representing a distinct aspect of hairstyle classification relevant to the booking context.

4.1 Label Dimensions

Named Styles (Key)

The most specific level of classification. Each named style corresponds to a distinct hairstyle with a recognised name within the domain. Multiple named styles may apply simultaneously to a single image when a style combines elements of more than one distinct technique.

Classes in this dimension:

alicia_keys, boho, box_braids, butterfly_locs, faux_locs, flat_twist, fulani_braids, invisible_locs, island_twist, lemonade_braids, marley_twists, microlocs, mini_twists, passion_twists, ponytails, puff, soft_locs, space_buns

Base Technique Categories (Base)

The installation technique underlying the named style. These are broader categories that group related named styles by the core technique used. A named style implies one or more base technique labels.

Classes in this dimension: braids, cornrows, crochet, locs, twist

Technique Modifiers (Other)

Structural characteristics of the installation that modify or further describe the base technique. These labels are independent of the named style and can apply in combination with any compatible base technique.

Classes in this dimension: box, crisscross, knotless, natural, stitch, stitch_braids

Size and Length (Size)

Dimensional attributes describing the scale of the individual braids, locs, or twists, and the overall length of the style. These attributes are annotated independently of technique.

Classes in this dimension: large, long, medium, short, small

Add-ons and Accessories (Addon)

Elements added to or incorporated into the base style during installation. These include accessories, styling choices, and material variations that are layered onto the base technique.

Classes in this dimension: barrel_ends, beads, curls, edges, extension, updo

4.2 Multi-Label Structure

Each image in the benchmark can carry labels from multiple dimensions simultaneously. A single image may be labelled with a specific named style, a base technique category, one or more technique modifiers, a size attribute, and one or more add-ons at the same time. All of these labels are simultaneously correct for that image.

This multi-label structure reflects the operational requirements of the booking context. Each label dimension carries independent information: technique category determines which professionals can deliver the service; named style determines the specific service listing; size affects duration and material quantities; add-ons affect pricing and product requirements. A complete and correct label set across all dimensions is what the system requires to route a booking accurately.

4.3 Relationship to Model Internal Taxonomy

The benchmark taxonomy is a subset of the labels Lig Vision was trained on. The model's internal taxonomy is broader, covering additional style categories across other regions and cultural contexts that are outside the scope of this benchmark. When the model is evaluated against the benchmark taxonomy, only the benchmark-relevant label dimensions are assessed. The model's full output is not evaluated here, and the benchmark results do not represent the model's complete capability.

5. Dataset

5.1 Overview

The benchmark dataset consists of 314 images of protective and textured hairstyles drawn from the Afro-Caribbean styling tradition. All 314 images are allocated to the test split. There is no train or validation split in this dataset; it is a pure evaluation set with no overlap with Lig Vision's training data.

The source dataset contains 320 images in total. The evaluation script filters to images with at least one valid ground truth label under the current named style or base technique taxonomy; the 6 excluded images have no such label and are therefore not scoreable.

The dataset is available at <https://universe.roboflow.com/testing-workspace-htshl/woman-hairstyles-bevyz> under a CC BY 4.0 licence.

5.2 Source and Construction

The dataset was assembled from two publicly available image sources: <https://universe.roboflow.com/hairstyles-ybypj/woman-hairstyles> and <https://universe.roboflow.com/traing/braids-detections>. Images were selected to provide representative coverage of the named styles included in the benchmark taxonomy.

Lig Vision has never been trained on any image from this dataset. The dataset was assembled after model training was complete, and images were not downloaded or processed locally during training. This is verified by the evaluation design: the benchmark code fetches images live from the public dataset URL rather than from a local copy, eliminating any possibility of data leakage. Both source datasets are publicly available on the web, and it is possible that images from those sources appear in the pre-training corpora of general-purpose vision-language models evaluated in this benchmark. This does not affect how those models' results are interpreted, as they are included as reference comparison points rather than as competing claims of zero-shot performance on strictly unseen data.

5.3 Labelling Methodology

Every image in the dataset was labelled by a professional hairstylist with active domain expertise. Labellers were not generalist annotators. This decision reflects the specific challenge of this domain: fine-grained distinction between related protective styles requires labellers who can recognise the relevant visual features in practice, not just identify them from descriptions.

Labels were applied across all five benchmark dimensions simultaneously for each image. Every image received a complete label set covering all applicable classes across named styles, base technique categories, technique modifiers, size attributes, and add-ons.

5.4 Verification

Every label set was independently verified by a second professional hairstylist. Disagreements were resolved by review. The dual-verification process is the mechanism by which label accuracy is assured. For a classification task where the correct label set requires domain expertise to produce, inter-annotator agreement between domain experts is the appropriate quality standard.

5.5 Label Distribution

The dataset is not uniformly distributed across all classes. Several named styles are well-represented with 50 or more images; others appear fewer than five times. This reflects the natural distribution of the styles represented in the source datasets and in professional hairstyle photography generally. Macro-averaged metrics are used throughout the evaluation, which gives equal weight to each class regardless of support. Per-class results are reported alongside macro averages to allow support-aware interpretation.

5.6 Known Label Convention Notes

The dataset was labelled under a standard professional annotation brief. Certain classes with very similar visual characteristics may have borderline cases where reasonable domain experts could disagree. These cases are most prevalent among closely related loc variants and among styles that share base techniques. Where ambiguity exists in the ground truth, it is reflected in lower per-class F1 scores for the affected classes and should be interpreted accordingly.

6. Comparison Models

The benchmark evaluates Lig Vision alongside three frontier general-purpose vision-language models. These models are not fine-tuned for hairstyle classification and receive no domain-specific training or prompting beyond the task instruction and the benchmark taxonomy. They are included to establish a reference baseline and to measure what specialised domain training contributes to classification performance.

6.1 GPT-5.5

OpenAI's GPT-5.5, queried via the Responses API with low reasoning effort and a maximum of 300 output tokens. The model receives the image and a structured prompt listing the benchmark taxonomy classes. It returns a JSON object with its predicted labels for each dimension.

6.2 Claude Opus 4.8

Anthropic's Claude Opus 4.8, queried via the Messages API with a maximum of 300 output tokens. The same structured prompt format is used as for GPT-5.5.

6.3 Gemini 2.5 Flash

Google's Gemini 2.5 Flash, queried via the Gen AI client with thinking disabled, temperature 0, and a maximum of 300 output tokens. The same structured prompt format is used.

6.4 Prompting Approach

All three frontier models receive the same prompt structure: the image, a task instruction to classify the hairstyle, and the valid label lists for each benchmark dimension. No few-shot examples are provided. No system prompt is used. The taxonomy is injected from `taxonomy.json` at evaluation time.

Frontier models return binary label lists (present or not present). They do not return probability scores. As a result, mean average precision (mAP) cannot be computed for frontier models. Only F1 is reported for these models.

7. Evaluation Methodology

7.1 Head-Split Evaluation

Each label dimension is evaluated independently on the subset of images that have at least one ground truth label for that dimension. Images with no ground truth for a given dimension are excluded from that dimension's evaluation.

This is essential for correctness. Including images with no ground truth for a given dimension in its evaluation silently inflates false positive counts for every class in that dimension. Any prediction the model makes for a dimension where the image has no ground truth is penalised as incorrect against an empty label set. Head-split evaluation removes this contamination.

Concretely: the named style F1 is computed on the 283 images that carry at least one named style label. The base technique F1 is computed on the 278 images with at least one base technique label. The sample sizes for other dimensions differ accordingly, and each is computed independently.

7.2 Metrics

Macro-averaged F1. For each class within a dimension, F1 is computed as the harmonic mean of precision and recall. Results are macro-averaged across all classes that have at least one positive ground truth example in the evaluation set. The macro average gives equal weight to each class regardless of support. Per-class F1 scores and support counts are reported alongside the macro average to allow support-aware interpretation.

Mean Average Precision (mAP). For Lig Vision, which returns probability scores rather than binary labels, mAP is computed as the macro-average of per-class average precision scores. Average precision measures the area under the precision-recall curve for a given class and is computed across all possible threshold values, not at any single threshold. mAP is a threshold-independent measure of ranking quality. A high mAP alongside a lower F1 typically indicates threshold miscalibration rather than poor discriminative ability. mAP is not reported for frontier models, which do not return probability scores.

7.3 Threshold Handling

Lig Vision returns a continuous probability score for each class. Whether a class is predicted as active depends on a per-class threshold applied to those probabilities.

Lig Vision is trained on a broader internal taxonomy than the benchmark covers. As a consequence of distributing probability mass across more classes than the benchmark taxonomy contains, raw probabilities for benchmark classes are structurally lower than they would be if the model were trained only on the benchmark taxonomy. Per-class threshold calibration corrects for this structural offset.

Per-class thresholds are calibrated using 5-fold cross-validation on the benchmark dataset. The dataset is partitioned into five equal folds. For each fold, thresholds are optimised on the remaining four folds by per-class F1 maximisation over a threshold grid; the held-out fold is then evaluated using those thresholds. The final threshold for each class is the mean of the five fold-specific optima. The F1 scores reported for Lig Vision are out-of-fold (OOF) metrics: each image is evaluated using thresholds calibrated on data that excludes that image. No image appears in its own calibration set. This ensures that reported F1 scores are honest estimates of generalisation, not estimates inflated by test-set contamination.

Classes with fewer than five positive examples within a calibration fold fall back to the model's native production threshold for that fold. The calibrated thresholds and calibration script are documented in the benchmark repository.

Frontier models produce binary outputs. No threshold is applicable.

7.4 Hierarchical Confidence Boosting

For Lig Vision, threshold relaxation is applied to named style classes when the corresponding base technique class exceeds a confidence threshold of 0.80. When a base technique class is predicted with high confidence, the thresholds for named style classes that structurally belong to that technique family are reduced by a factor of 0.65, subject to a minimum floor of 0.20. This reflects the informational relationship between base technique and named style: strong base technique confidence is evidence that should modestly lower the decision threshold for compatible named styles. Classes whose threshold is already at or below the floor do not receive the boost. The boosting parameters — the confidence threshold of 0.80, the reduction factor of 0.65, and the floor of 0.20 — are fixed constants established during model development and set in production prior to this benchmark. They were not fitted, searched, or tuned using benchmark data. The leakage-free guarantee described in Section 7.3 covers the per-class decision thresholds; this guarantee extends equally to the boosting parameters, which are production-fixed and benchmark-independent.

7.5 Reproducibility

The benchmark is fully reproducible by external parties with the following provisions:

- Frontier model results require API credentials for OpenAI, Anthropic, and Google. These are provided by the tester and are not included in the repository.
- Lig Vision results are pre-computed and included in `results/lig\_vision.json`. The Lig Vision inference endpoint is not published. Individual predictions can be verified via the Ligax AI playground at `https://www.ligaxai.com/playground`.
- The dataset is fetched live from `https://universe.roboflow.com/testing-workspace-htshl/woman-hairstyles-bevyz` at evaluation time. No local download is required.
- All evaluation code is in the benchmark repository at `https://github.com/Ligax-LTD/Lig\_benchmark\_test`.

8. Current Results

Results reflect the benchmark evaluation run on 14 June 2026. All metrics use the head-split evaluation protocol. Lig Vision named style F1 and base technique F1 are out-of-fold estimates from 5-fold cross-validated threshold calibration (see Section 7.3). mAP is threshold-independent and is computed within the same head-split evaluation subset as F1.

8.1 Named Styles (Key) and Base Technique Categories

Model	Named Styles F1	Named Styles mAP	Base Technique F1
Lig Vision	0.622	0.770	0.780
Claude Opus 4.8	0.477	n/a	0.693
Gemini 2.5 Flash	0.433	n/a	0.667
GPT-5.5	0.411	n/a	0.572

Sample sizes: 283 images evaluated for named styles; 278 images for base technique categories.

8.2 Named Styles -- Per-Class Detail (Lig Vision)

F1 scores are out-of-fold. mAP is threshold-independent.

Class	F1	mAP	Support
alicia_keys	1.000	1.000	2
boho	0.886	0.944	79
box_braids	0.716	0.859	52
butterfly_locs	0.444	0.626	7
faux_locs	0.784	0.752	43
flat_twist	0.857	1.000	3
fulani_braids	0.828	0.886	14
invisible_locs	0.000	0.317	2
island_twist	0.167	0.432	5
lemonade_braids	0.286	1.000	1
marley_twists	0.817	0.870	61
mini_twists	0.773	0.912	21
passion_twists	0.700	0.681	9
ponytails	0.783	0.921	12
puff	0.821	0.963	18
soft_locs	0.316	0.228	17
space_buns	0.400	0.700	2

Note: microlocs is included in the named style taxonomy but has no effective ground truth in the current dataset version due to a label naming discrepancy (dataset uses `micro\_locs`; this is resolved in the dataset for the next benchmark version).

8.3 Base Technique Categories -- Per-Class Detail (Lig Vision)

Class	F1	mAP	Support
braids	0.932	0.967	94
cornrows	0.500	0.598	5
crochet	0.667	0.698	6
locs	0.863	0.878	75
twist	0.937	0.951	109

8.4 Additional Dimensions

Results for technique modifiers, size and length, and add-ons are pending the full-taxonomy evaluation run and will be published in the next benchmark update.

9. Reproducing Results

9.1 Requirements

Python 3.10+
requests
python-dotenv
numpy
scikit-learn
openai
anthropic
google-generativeai

Install dependencies:

```
pip install -r requirements.txt
```

9.2 Configuration

Create a `.env` file in the benchmark root:

```
RF_API_KEY=<your Roboflow API key>  
RF_WORKSPACE=testing-workspace-htshl  
RF_PROJECT=woman-hairstyles-bevyz  
RF_VERSION=1
```

```
OPENAI_API_KEY=<your OpenAI API key>  
OPENAI_MODEL=gpt-5.5
```

```
ANTHROPIC_API_KEY=<your Anthropic API key>  
ANTHROPIC_MODEL=claude-opus-4-8
```

```
GOOGLE_API_KEY=<your Google API key>  
GOOGLE_MODEL=gemini-2.5-flash
```

A Roboflow API key is required to fetch the dataset. This can be obtained from Roboflow account settings at `app.roboflow.com/settings/api`. A free Roboflow account is sufficient. The frontier model API keys are obtained from the respective provider accounts.

The `LIG\_VISION\_URL` variable is intentionally omitted. Lig Vision results are pre-computed and committed to `results/lig\_vision.json`. Running the benchmark without this variable will skip the Lig Vision model and use the cached results for comparison. Individual Lig Vision predictions can be verified at <https://www.ligaxai.com/playground>.

9.3 Running the Benchmark

To evaluate all frontier models:

```
python evaluate.py
```

To evaluate a specific model:

```
python evaluate.py --models gpt
python evaluate.py --models claude
python evaluate.py --models gemini
```

To calibrate Lig Vision thresholds and update `results/lig\_vision.json` with out-of-fold metrics:

```
python calibrate.py
```

Results are written to `results/<model>.json` and include per-image predictions, per-class metrics, and summary statistics.

9.4 Verifying Lig Vision Results

The Ligax AI playground at <https://www.ligaxai.com/playground> accepts any image and returns the same model output that produced the benchmark results. To verify a specific benchmark prediction:

1. Open the dataset at <https://universe.roboflow.com/testing-workspace-htshl/woman-hairstyles-bevyz>
2. Download any image from the test split
3. Upload it to the playground
4. Compare the returned labels with the corresponding entry in `results/lig\_vision.json`

The playground uses the same inference endpoint and the same model version as the benchmark. Results should be consistent with the cached benchmark output for images that fall within the benchmark domain (Afro-Caribbean protective hairstyles).

10. Limitations and Known Issues

Low-support classes. Several named style classes have very few ground truth examples in the current dataset (fewer than five images). F1 scores for these classes have high variance and should be interpreted with caution. Classes with support of two or fewer are particularly unreliable as benchmark indicators. This is documented in the per-class results table.

Microlocs naming discrepancy. The `microlocs` class in the benchmark taxonomy does not match the `micro\_locs` column in the current dataset due to a labelling normalisation mismatch. This class has effectively zero ground truth in the current dataset version and is excluded from macro F1 computation. It will be resolved in the next dataset version.

Frontier model threshold. Frontier models return binary labels without probability scores. The decision threshold is implicit in the model's output and cannot be calibrated post-hoc. This means frontier model F1 scores reflect both classification quality and the models' implicit confidence thresholds, which may not be optimal for the benchmark taxonomy.

Convention gap for base technique labels. The benchmark dataset was labelled under a standard annotation convention. Lig Vision was trained with an enriched label convention in which certain named styles and modifier combinations imply specific base technique labels that annotators may not have recorded. This means the model may fire base technique labels on images where the benchmark ground truth does not record them, producing false positives that are domain-correct but benchmark-incorrect. Base technique thresholds are calibrated to compensate for this effect.

Benchmark scope. This benchmark covers one domain (Afro-Caribbean protective hairstyles) and one regional style tradition. Lig Vision is capable across a broader range of styles and demographics. The benchmark results are not representative of the model's full performance envelope.

11. Ongoing Development

The Ligax AI hairstyle classification benchmark is under active development. The following extensions are planned or in progress.

Full-taxonomy evaluation. The current published results cover named styles and base technique categories. Results for technique modifiers, size and length attributes, and add-ons and accessories will be published in the next benchmark update following the completion of internal evaluation across all label dimensions.

Dataset expansion. The current dataset covers 314 images. Subsequent versions will increase coverage across all named style classes, with particular attention to classes that currently have low support. New images will be subject to the same dual-verification labelling process as the original dataset.

Additional style domains. Future benchmark versions will extend to hairstyle categories beyond the Afro-Caribbean protective hairstyle domain covered here. This will include styling traditions from other regions and demographic contexts in which Lig Vision is already trained and deployed.

Continuous model versioning. Each model version is evaluated against the current benchmark version at time of release, and results are published with both the model version and benchmark version clearly paired.

Model additions. The comparison set of frontier models will be updated as new vision-language models are released and as model versions used in the initial benchmark are superseded.

12. Version History

Version	Date	Changes
1.0	14 June 2026	Initial release. Named styles and base technique categories. 314 images. Lig Vision, GPT-5.5, Claude Opus 4.8, Gemini 2.5 Flash.